

Bike Rental Demand Forecasting

From Data Preparation to Model Comparison and Critical Reflection

Anton Wilhelmi and Alexander Harant
Introduction to Machine Learning
Group Project - Deliverable 2

Abstract—This report documents the full workflow after the first problem formulation and integrates the earlier data preparation summary into one continuous methodological story. We explain how the raw bike rental data was reduced to a clean modelling table, why each transformation was made, how the regression pipeline was built, which models were trained, how they were evaluated, and what the current results mean. The final reduced dataset contains one target, `total_rentals`, and 17 numeric predictors. We then compared eight regression models with the same chronological train-validation-test split and the same evaluation metrics. The comparison shows that ensemble tree methods currently perform best: Gradient Boosting is the best validation model, while Random Forest is the best model on the held-out test set. At the same time, the results also show clear weaknesses in the current representation, especially the limited treatment of station identity and time structure. We therefore end the report with a critical discussion of what still limits the system and what should be improved next.

I. INTRODUCTION

The goal of the project is to predict bike rental demand with machine learning in a way that is technically sound, reproducible, and easy to justify.

This report therefore combines two pieces of work into one coherent document. First, it includes the complete reasoning behind the final reduced dataset, because the modelling results only make sense if the data representation is clear. Second, it describes the modelling pipeline that starts from this reduced dataset and compares several regression models under the same evaluation protocol.

II. FROM DELIVERABLE 1 TO DELIVERABLE 2

After the feedback on Deliverable 1, we decided to change the dataset. One important reason was that the new dataset contains much more data than the old one. This gives a stronger basis for model training, testing, and comparison. A larger dataset is useful because the models can learn from more examples and the final results become more reliable. For this reason, we switched to the Capital Bike-share dataset from Kaggle: <https://www.kaggle.com/datasets/taweilo/capital-bikeshare-dataset-202005202408>

This change was important because the new dataset fits the task more directly. Our goal is to predict bike rental demand. The new dataset provides a good basis for this because it

includes rental data, station information, and weather information in a form that can be transformed into a daily demand prediction dataset.

The first deliverable already defined the task as a regression problem. It also identified important core ideas: the target should be a rental count, time-aware splitting is necessary, and leakage variables must be removed. These ideas stayed central in Deliverable 2.

However, after Deliverable 1 we made several important adjustments. The main change is that we built a much clearer data processing pipeline. Instead of moving directly from raw data to modelling, we first created a structured daily station-level dataset, analysed feature redundancy, simplified the target, and built one final numeric modelling table. This step was necessary because a clean and transparent input table is the basis for a fair comparison of models.

We also expanded the evaluation. In Deliverable 2 we used MAE, RMSE, Median Absolute Error, MAPE, R^2 , and Explained Variance. This gives a broader and more reliable picture of model behaviour.

A further change is the modelling structure itself. We created one shared project structure for all models. Every model now uses the same utilities, the same split logic, the same target variable, and the same output format. This makes the experiments directly comparable. It also reduces the risk that one model looks better only because it was trained or evaluated under a different setup.

III. DATA PREPARATION METHODOLOGY

A. Why we reduced the raw data

The raw data was delivered in four files contained detailed rental data, station information, usage frequency data, and weather data. Working with the raw tables directly would have made the modelling phase harder to interpret and more error-prone. Therefore, the first goal was to build one clean analytical table that could later be used for regression.

We decided to focus on the 20 most used start stations. This was a practical and methodological decision. Practically, it reduced the size of the problem (2,8 GB file) and made early experimentation faster. Methodologically, it also reduced noise from stations with very few observations. Rare stations can introduce unstable patterns, while highly used stations provide a stronger and more consistent signal. The idea was not to

claim that rare stations are unimportant, but to create a stable first forecasting setup before scaling the problem further.

The next step was to aggregate the rental information at daily station level and merge it with weather variables. This aggregation created a table where every row corresponds to one station-day combination. This step was important because it aligned the demand outcome with the explanatory variables we wanted to use. If weather is available daily, then a daily demand representation is also a natural first modelling level.

B. Why we stored station names separately

At one stage the processed data still contained both `start_station_id` and station names. We removed the name columns from the modelling table and stored the mapping in a separate file. The reason is simple: station names are descriptive labels for humans, but they do not add predictive information beyond the station identifier. Keeping both inside the modelling matrix would only duplicate the same information and make the dataset less clean.

This decision also made later steps easier. The modelling table stays compact and numeric, while the mapping file still allows us to interpret results by station when needed.

C. Why we changed the date representation

The raw date column was not kept in its original form in the final modelling file. A string-based date is not directly useful for the regression models we trained. Instead of leaving the raw date string in the table, we represented time through numeric variables such as month, weekday, year and a chronological index called `time_idx`.

This decision had two advantages. First, the final input matrix became directly trainable because all predictors were numeric. Second, we did not lose the time structure of the data. The new variables still preserve seasonality and order, but in a format that the models can process.

D. Why we analysed correlations before modelling

Before creating the final reduced dataset, we computed a correlation matrix and a heatmap for the predictors. The purpose was to identify variables that carried almost the same information. This matters because redundant predictors can make models less stable and harder to interpret. The issue is especially relevant for linear models, where strong collinearity can distort coefficients.

The analysis showed that several weather variables were nearly duplicates. Examples included:

`solarenergy` vs. `solarradiation`,
`feelslike` vs. `temp`,
`feelslikemin` vs. `tempmin`,
`feelslikemax` vs. `tempmax`.

We kept the more interpretable or slightly more informative variable and removed the weaker duplicate. This reduced unnecessary redundancy and made the final dataset easier to explain. We did not try to remove every possible relationship between predictors, because some correlation is natural in weather data, but we removed the strongest near-duplicates that offered almost no new information.

E. Why we simplified the target to one output

During preprocessing, the dataset still contained multiple rental-related count columns, for example counts for members and casual users or counts for specific bike types. At first glance these columns look useful, but they create a serious modelling problem when the main prediction target is `total_rentals`. The reason is that these sub-counts are structurally tied to the total and may allow the model to reconstruct the answer almost directly.

To avoid this issue, we removed the count decomposition and kept only one target variable:

`total_rentals`

This was a very important decision. It turned the task into one clean regression problem and avoided a leakage-like situation in which the model would partly learn the target from variables that are conceptually too close to it.

F. Why we did not apply PCA

We also considered whether a PCA-based representation would be useful. We decided against it. After the correlation-based reduction, the number of predictors was already manageable. PCA would have transformed the remaining features into mixtures of original variables, which would make the model harder to interpret. Since this project strongly values explainability and critical reflection, preserving the meaning of the original predictors was more valuable than compressing them further.

G. Final reduced dataset

The final modelling table contains one target and 17 predictors. The predictors are:

`start_station_id`, `tempmax`, `humidity`,
`precip`, `precipcover`, `cloudcover`, `windspeed`,
`visibility`, `sealevelpressure`, `uvindex`,
`sunset_minutes`, `month`, `weekday`, `year`,
`time_idx`, `snow`, `snowdepth`

This final dataset is compact, fully numeric, and directly usable for machine learning. It is also the direct result of several explicit design decisions rather than an arbitrary column selection.

IV. MODELLING PIPELINE AFTER DATA PREPARATION

A. Why we created a shared modelling structure

Once the final dataset was ready, we built a structured modelling folder. Each model has its own subfolder, and all common functionality is stored in a `common` folder. This decision was made for reproducibility, consistency, and easier debugging.

The `common` folder contains the shared code for paths and settings, preprocessing, splitting, metrics, plotting, and file saving. This means that every model script follows the same core logic: load the same reduced dataset, build the same time-aware split, predict the same target, compute the same metrics, and save outputs in the same format. This matters because a model comparison is only fair if the evaluation procedure is consistent.

B. Why we used a chronological split

We used a chronological split with 70% training data, 15% validation data, and 15% test data. We did not use a random split. This is one of the most important methodological decisions in the report.

The reason is that bike demand is a time-based problem. In a real forecasting scenario, a model is trained on older data and then asked to predict future demand. A random split would mix older and newer observations and would therefore make the problem unrealistically easy. It could also create information leakage across time. The chronological split respects the natural order of the observations and produces a more honest estimate of future performance.

The variable `time_idx` was used to maintain the exact order of the station-day observations. Because the split is stricter than a random split, the resulting scores are harder to optimize, but they are also much more realistic.

C. Why we scaled only selected models

Feature scaling was applied only where it is methodologically appropriate. Linear models, regularized linear models, and KNN are sensitive to feature magnitude, so scaling is necessary there. Otherwise a feature with a large numeric range can dominate other predictors for purely numerical reasons.

Tree-based models work differently. They split on thresholds and do not depend on Euclidean distance or coefficient penalty in the same way. Therefore, scaling is not required for Decision Tree, Random Forest, or Gradient Boosting. This selective scaling choice keeps the preprocessing aligned with the modelling assumptions.

V. MODEL SELECTION STRATEGY AND JUSTIFICATION

We deliberately started with simple models and then moved towards more flexible models. This follows the lecture logic of establishing a baseline, then adding complexity only when it is justified.

A. Dummy Regressor

The Dummy Regressor was used as the minimum reference point. It does not learn real relationships from the data. Its role is to answer a simple question: do the real models extract useful signal at all? If a real model cannot improve over the dummy baseline, then either the representation is too weak or the modelling setup is flawed.

B. Linear Regression

Linear Regression was the first real predictive model because it is easy to interpret and acts as a clean benchmark. It tells us whether the final reduced features already contain a meaningful approximately linear signal. It is also a good diagnostic model because poor performance here can reveal weak representation or strong nonlinearity.

C. Ridge Regression

Ridge Regression extends Linear Regression by shrinking coefficients. This is useful when predictors are still somewhat correlated, which is likely in weather and calendar data. The reason for including Ridge was therefore not only performance but also stability. If Ridge improves over plain Linear Regression, that is evidence that mild regularization is helpful.

D. Lasso Regression

Lasso was included because it regularizes even more aggressively and can shrink some coefficients toward zero. This makes it useful when we suspect that not all predictors are equally helpful. Lasso therefore adds a feature-selection perspective on top of regularized regression.

E. Decision Tree

A Decision Tree was tested because it can model nonlinear interactions and threshold effects without requiring manual feature transformations. This is appealing for weather-driven demand, where effects are often not purely linear. However, single trees can overfit, so they also serve as a useful contrast between interpretability and stability.

F. KNN Regressor

KNN was included as a local neighbourhood model. The idea behind it is simple: similar observations may have similar target values. KNN can capture local structure that linear models miss, but it is highly sensitive to scaling and often becomes slower at prediction time. Including it helped us compare a distance-based strategy with coefficient-based and tree-based strategies.

G. Random Forest

Random Forest was a natural next step after Decision Tree. A single tree is flexible but unstable. Random Forest reduces this instability by combining many trees trained on bootstrapped samples. It is therefore a good candidate when a single tree seems promising but too variable. In course terms, it is also a bagging method that can reduce variance.

H. Gradient Boosting

Gradient Boosting was included because boosting can reduce bias by building trees sequentially and focusing on the remaining errors. This makes it a strong candidate when simpler models underfit the demand structure. It is typically more powerful than a single tree, but also more computationally expensive and more sensitive to hyperparameter choices.

VI. EVALUATION STRATEGY AND METRICS

We used several regression metrics because one single number cannot describe model quality well enough. Each metric shows a different side of the prediction error. This is important in our bike rental problem, because we do not only want a model that is good on average. We also want to understand whether it makes large mistakes, whether it performs consistently, and whether it explains the main demand patterns.

Mean Absolute Error (MAE) measures the average absolute difference between the predicted rentals and the real rentals. It tells us, in simple units, how many bikes the model is off on average. For example, an MAE of 20 means that the prediction is wrong by about 20 rentals on a typical day. This metric is very easy to understand, because it is expressed in the original target unit. For our problem, MAE is relevant because the final goal is to predict the number of rented bikes as directly and clearly as possible.

Root Mean Squared Error (RMSE) stands for Root Mean Squared Error. It is also measured in rental units, but it punishes large errors more strongly than MAE. This happens because the errors are squared before they are averaged. As a result, a few very large mistakes increase RMSE a lot. This is useful in our problem because big misses are especially problematic. If the model strongly underestimates or overestimates bike demand on a day with unusual weather or strong seasonal effects, this can be much more harmful than many small mistakes.

Median Absolute Error measures the median of the absolute prediction errors. In simple words, it tells us what a typical middle-sized error looks like. It is more robust than MAE and RMSE because it is less influenced by a few extreme cases. This is relevant for our dataset because bike demand can vary strongly across stations and days. Some days may have unusual peaks or drops. Median Absolute Error helps us see how the model performs in the more normal cases, without being dominated by rare outliers.

Mean Absolute Percentage Error (MAPE) stands for Mean Absolute Percentage Error. It measures the prediction error relative to the real value and expresses it as a percentage. This helps us understand whether the model is wrong by a small or large proportion of the actual demand. For example, an error of 10 rentals is much more serious when the true demand is 20 than when it is 100. That is why MAPE gives an additional perspective that MAE and RMSE do not show. For our problem, this is useful because rental counts are not equally large on all days. However, MAPE must be interpreted carefully when the true number of rentals is very small, because percentage errors can then become unstable.

R^2 , also called the coefficient of determination, measures how much of the variation in bike rentals is explained by the model. A higher R^2 means that the model captures more of the underlying demand structure. If R^2 is close to 1, the model explains the data very well. If it is close to 0, the model is not much better than a simple average prediction. If it is negative, the model performs even worse than this baseline. In our case, R^2 is relevant because bike rentals are influenced by many systematic factors such as weather, time, and station effects. We therefore want to know how much of this general pattern the model can actually learn.

Explained Variance measures how much of the variance of the target is captured by the predictions. It is similar to R^2 , but it focuses more directly on how well the spread of the target values is represented. This gives another view on the same general question: does the model only predict

around the mean, or does it also capture the ups and downs of real bike demand? For our problem, this is important because daily rentals do not stay constant. Good models should reflect changes in demand across time and weather conditions, not only produce flat average-like predictions.

Using all of these metrics together gives a more complete and honest evaluation. A model may have a low MAE but still make some very large mistakes, which would be visible in RMSE. Another model may explain more variance but still have weaker percentage-based performance, which would appear in MAPE. Looking at several metrics at the same time helps us compare models more carefully and choose a model that is not only strong in one narrow sense, but useful for the full prediction task.

VII. RESULTS AND MODEL COMPARISON

Table I summarizes the most relevant validation and test results for all eight models.

The first clear result is that all real models outperform the dummy baseline by a large margin. This is an important positive finding because it shows that the reduced dataset contains real predictive signal.

The second result is that the linear family improves in a meaningful sequence: Linear Regression is clearly better than Dummy, Ridge improves slightly over Linear, and Lasso improves further. This tells us that regularization helps, but it also suggests that the demand problem is not purely linear.

The third and strongest result is that the ensemble tree models are currently the best performers. Gradient Boosting achieves the best validation RMSE (27.00) and the best validation R^2 (0.474), while Random Forest achieves the best held-out test RMSE (30.31), the best test MAE (22.87), and the best test R^2 (0.406). This means that the two strongest models in the current project phase are the two ensemble tree methods, with Random Forest currently being the safer final choice because the test set is the most important final evaluation set.

TABLE I: Main comparison results for all evaluated models

Model	Val. MAE	Val. RMSE	Val. R^2	Test MAE	Test RMSE	Test R^2	Fit Time (s)	Predict Time (s)
Dummy	37.00	46.10	-0.533	38.14	48.56	-0.523	0.001	0.000
Linear	26.14	32.94	0.217	29.40	36.72	0.129	0.018	0.002
Ridge	25.57	32.37	0.244	28.68	35.91	0.167	0.003	0.000
Lasso	24.04	31.13	0.301	26.21	33.49	0.275	0.027	0.002
Decision Tree	21.33	28.64	0.408	24.32	32.03	0.337	0.109	0.005
KNN	23.50	31.14	0.301	24.42	32.35	0.324	0.001	0.832
Random Forest	20.40	27.28	0.463	22.87	30.31	0.406	0.872	0.185
Gradient Boosting	20.53	27.00	0.474	23.71	30.57	0.396	6.232	0.050




Fig. 1: Overview of the most important comparison dimensions. The figure highlights that Gradient Boosting is strongest on validation, while Random Forest is strongest on the held-out test set.

VIII. VISUAL COMPARISON OF THE MODELS




Fig. 2: RMSE and MAE comparison across train, validation, and test splits. These plots show that ensemble trees reduce error substantially compared with the baseline and the linear models.




Fig. 3: R^2 and explained variance comparison. The same pattern appears in both views: the ensemble models explain the largest share of the target variation.




Fig. 4: MAPE and median absolute error. These metrics support the same overall conclusion, while also showing that some models differ in how stable their typical errors are.

IX. INTERPRETATION OF THE CURRENT RESULTS

The model comparison leads to several important conclusions.

First, the feature engineering work clearly mattered. The large gap between Dummy and the real models shows that the final reduced dataset is not arbitrary. It contains useful signal from weather, time, and station-related information.

Second, the improvement path across model families is informative. The move from Linear to Ridge and then to Lasso shows that regularization is helpful, but the gains stay




Fig. 5: Ranking by validation RMSE and test RMSE. Gradient Boosting wins the validation phase, but Random Forest becomes the best model on the final test set.

limited. This suggests that the main remaining challenge is not simply coefficient instability. The problem likely contains nonlinear interactions and station-specific effects that linear models cannot capture well.

Third, the tree-based models confirm this interpretation. Decision Tree already improves strongly over the linear family, and the ensemble versions improve even more. This strongly suggests that the demand process includes nonlinear behaviour. Examples that fit this pattern are threshold effects in weather, different station sensitivities, and interactions between season, daylight, and rain.

Fourth, the current comparison also reveals that model selection depends on the evaluation stage. Gradient Boosting is best on validation, which means it was the strongest candidate during development. Random Forest is best on the final held-out test set, which makes it the most convincing model at the current stage if one model must be selected for reporting. This is an important distinction: the best development score and the best final score do not have to be the same.

X. DIRECT WEAKNESSES VISIBLE IN THE CURRENT SETUP

Even though the results are promising, the current system still has important weaknesses. These weaknesses are not small details; they directly explain why the performance is still moderate and why the report should remain critical.

The first weakness is the representation of `start_station_id`. In the final reduced table it is still treated as a numeric variable. However, a station identifier is




Fig. 6: Rank-sum comparison across multiple metrics. This gives a broader summary than a single metric and again shows Gradient Boosting on top for validation and Random Forest on top for the final test comparison.

not a true numeric quantity. There is no meaningful distance between station IDs, and a linear model may interpret the values in a way that does not match reality. This is especially problematic for the linear family. The current representation is practical, but not ideal.

The second weakness is the simplified daily aggregation. Daily modelling is easier to manage and useful for a first pipeline, but it removes many short-term demand patterns. Bike rental demand often depends on commuting peaks, hour-specific weather changes, or short event windows. Aggregating everything at the day level hides that structure.

The third weakness is the missing lag information. The current predictors do not explicitly include yesterday's rentals, last week's rentals, or rolling averages. In a forecasting problem this is a major limitation, because rental demand often has temporal dependence. Without lag features, the model only sees contemporaneous weather and calendar information and cannot directly use short-term memory.

The fourth weakness is the limited context coverage. The current dataset does not explicitly include holidays, public events, road works, local disruptions, transport strikes, or unusual city-specific situations in the final modelling file. This means that some strong demand changes remain unexplained. The model is then forced to treat them as noise.

The fifth weakness is that the problem is evaluated in a realistic chronological way. This is methodologically correct, but it also means the future period may differ from the earlier training period. In other words, concept drift is a real issue

here. A model that fits the earlier data well may still struggle if later behaviour changes.

The sixth weakness is that some models already show a clear train-validation gap. KNN is the strongest example: the training fit is extremely strong, but the validation and test results are much worse. This is a typical sign of overfitting or overly local memorization. Decision Tree also shows the same basic pattern, though less extremely. The linear family shows the opposite tendency: it is safer but also more limited, which suggests underfitting relative to the more flexible ensemble methods.

XI. LONG CRITICAL REFLECTION AND NEXT STEPS

At this stage, the project has a strong and reproducible structure, but it is not yet a finished forecasting solution. The most important insight is that adding models alone is not the main answer. The comparison already includes several fundamentally different approaches: a trivial baseline, linear regression, regularized regression, a single tree, a local neighbourhood method, and two ensemble tree methods. Because the strongest gains appear when moving to models that can capture nonlinear structure, we now know that the task is more complex than a simple linear relationship between weather and total rentals. However, the test scores also show that even the best model still leaves a large part of the variance unexplained. That means the next improvement should probably come more from better representation than from adding yet another algorithm without changing the data.

The first improvement should be a better encoding of station identity. A raw station ID is a convenient placeholder, but not a meaningful numeric signal. A one-hot representation, target-aware categorical strategy, or even station-specific models would be more appropriate. The second improvement should be richer temporal features. Month and weekday are useful, but they still represent time in a rather simple way. Cyclical encodings, lag variables, rolling means, and demand history from previous days or weeks would likely provide much more predictive structure. The third improvement should be a closer residual analysis. The current aggregated plots tell us which model is better overall, but they do not yet tell us where the models fail most strongly. Residuals should be analysed by station, season, weather regime, and demand level. This can show whether the model fails on high-demand peaks, rainy days, winter periods, or a subset of stations.

A fourth next step is to add learning-curve style diagnostics. The current results already hint at both phenomena in different places: the linear models appear safer but limited, while KNN and some tree models fit the training data much more strongly than the future data. A proper learning-curve analysis would make this diagnosis more explicit. A fifth next step is to reconsider whether daily aggregation is the right long-term representation. It was a good decision for building the first stable pipeline, but hourly modelling may better match the original demand process if the project later aims for more precise operational forecasting.

We transformed a raw data problem into a reproducible machine learning pipeline with clear modelling decisions and a transparent comparison of multiple approaches. The best current final model is Random Forest on the held-out test set, while Gradient Boosting was the strongest development model on validation. That is already a meaningful result. At the same time, it makes it equally clear that the most valuable next progress will likely come from improving the data representation, adding temporal memory, and analysing model failures more deeply rather than only increasing algorithmic complexity.

XII. PRELIMINARY ETHICAL CONSIDERATIONS

Predicting bike rental demand is generally a low-risk machine learning task. However, we still need to consider some ethical and social effects.

First, we decided to use only the 20 most used start stations. This was necessary to build a stable first model, but it creates a representation bias. By leaving out smaller stations, our model ignores the demand in areas that might already have fewer services. In a real-world system, this could lead to an unfair distribution of bikes, where less popular neighborhoods get even less attention.

Furthermore, the data we use is completely anonymous and grouped by day and station. This means no personal user information is exposed. If we change our model in the future to use hourly data or specific user types (like registered members or casual users), we must keep this strict privacy standard.

In future steps of this project, we should also think about the social and economic background of the stations. We must make sure that our model does not systematically disadvantage certain parts of the city.

XIII. AI TOOL USE

AI tools were used for drafting support, restructuring text, and helping to standardize code and reporting outputs. The final content, decisions, code execution, result interpretation, and critical reflection were reviewed and organized by the authors.